

SYSTEMS DEVELOPMENT FOR COMPUTATIONAL SCIENCE

LECTURE 11

Fabian Wermelinger

Harvard University

CS107 / AC207

Thursday, October 6th 2022

LAST TIME

- Python packages and the [Python Package Index](#)
- Towards automatic differentiation
- Linearization of Euler equations as an example for the Jacobian
- Newton's method

TODAY

Main topics: *Chain rule, Gradient operator and directional derivative, Automatic differentiation: forward mode*

Details:

- Example of Newton's method and numerical schemes to approximate derivatives.
- Evaluation trace
- The computational graph
- Computing derivatives of one variable using the forward mode


NEWTON'S METHOD: INTERSECTION OF TWO LINES

Given two functions $y_1 = x$ and $y_2 = \exp(-2(\sin(4x))^2)$.

Find x such that $y_1 = y_2$. The statement is equivalent to:



$$f(x) = x - \exp(-2(\sin(4x))^2) = 0.$$

Real world application: *ray-tracing*



NEWTON'S METHOD: INTERSECTION OF TWO LINES

Before we start, it is a good idea to visualize the problem:



There are *three* zeros and we can not solve this problem by hand.
Method of choice → Newton's method

NEWTON'S METHOD: INTERSECTION OF TWO LINES

How should the algorithm in our program look like?

Algorithm sketch:

- Need an *initial guess*
- Need some *termination criterion*
- Want to protect from infinite iterations if there is *divergence*
- *Optional:* Would like to pass *parameters as arguments*

```
1 x_k # initial guess
2 tol # convergence tolerance
3 max_it # maximum iterations
4 for k in range(max_it): # iteration loop
5     dx_k = -f(x_k) / dfdx(x_k) # compute correction
6     if abs(dx_k) < tol: # check for convergence
7         root = x_k + dx_k
8         break
9     x_k += dx_k # update the iteration variable
```

NEWTON'S METHOD: INTERSECTION OF TWO LINES

We need to determine the Jacobian J of $f(x)$:

Given the function

$$f(x) = x - \exp\left(-2(\sin(4x))^2\right),$$

find the derivative df/dx (*for this scalar case $J(x) = df/dx$*).

Options:

1. On a piece of paper
2. Using software to calculate derivatives *analytically*, e.g. [Mathematica](#).
Python also supports this with the [SymPy](#) package for symbolic math. You can install the package with

```
1 $ python -m pip install sympy
```

See the [online documentation](#) for more. *You can find a Python script and a Jupyter notebook for this example in the lecture code handout.*

NEWTON'S METHOD: INTERSECTION OF TWO LINES

Once the Jacobian $J(x)$ is known we can complete Newton module:

```
1 import numpy as np
2
3 f = lambda x: x - np.exp(-2.0 * np.sin(4.0 * x) * np.sin(4.0 * x))
4 J = lambda x: 1.0 + 16.0 * np.exp(-2.0 * np.sin(4.0 * x)**2
5               ) * np.sin(4.0 * x) * np.cos(4.0 * x)
6
7 def newton(f, J, x_k, tol=1.0e-8, max_it=100):
8     root = None
9     for k in range(max_it):
10         dx_k = -f(x_k) / J(x_k)
11         if abs(dx_k) < tol:
12             root = x_k + dx_k
13             print(f"Found root {root:e} at iteration {k+1}")
14             break
15         print(f"Iteration {k+1}: Delta x = {dx_k:e}")
16         x_k += dx_k
```

(To handle arguments we can use the [argparse](#) module from the Python standard library)

You can find this code in the *newton.py* module in the class repository.

NEWTON'S METHOD: INTERSECTION OF TWO LINES


Recall: the check whether `__name__` corresponds to the `'__main__'` string in the top-level scope *allows us to run our module just like a program!*

```
1 $ chmod 755 newton.py
2 $ ./newton.py --help
3 usage: newton.py [-h] -g INITIAL_GUESS [-t TOLERANCE] [-i MAXIMUM_ITERATIONS]
4
5 Newton-Raphson Method
6
7 optional arguments:
8   -h, --help            show this help message and exit
9   -g INITIAL_GUESS, --initial_guess INITIAL_GUESS
10                        Initial guess
11   -t TOLERANCE, --tolerance TOLERANCE
12                        Convergence tolerance
13   -i MAXIMUM_ITERATIONS, --maximum_iterations MAXIMUM_ITERATIONS
14                        Maximum iterations
15 $ ./newton.py --initial_guess 0.1
16 Iteration 1: Delta x = 1.218877e-01
```

→ for the initial guess 0.1, the method finds *one* root at 2.473652e-01

NEWTON'S METHOD: INTERSECTION OF TWO LINES

Validation of the result:



```
1 >>> from newton import (f, J, newton)
2 >>> root = newton(f, J, 0.1)
3 Iteration 1: Delta x = 1.218877e-01
4 Iteration 2: Delta x = 2.339599e-02
5 Iteration 3: Delta x = 2.066548e-03
6 Iteration 4: Delta x = 1.500080e-05
7 Found root 2.473652e-01 at iteration 5
8 >>> f(root)
9 5.551115123125783e-17 # about zero
```

Note that the initial guess is crucial:

```
1 >>> newton(f, J, 0.6)
2 Found root 6.692328e-01 at iteration 5 # converges to second root!
3 >>> newton(f, J, 0.9)
4 Found root 8.560317e-01 at iteration 4 # converges to third root!
```

NEWTON'S METHOD: INTERSECTION OF TWO LINES

Summary:

- We worked through a root finding problem to solve for the intersection of two lines based on a nonlinear function. An example application where this problem arises is ray-casting.
- Finding the solution to this problem required use of Newton's method which requires knowledge of the *derivative* of the system governing function $f(x)$.
- For a general setup in 3D space, there is not just one of these derivatives! In this case we call these "derivatives" the *Jacobian* of the system. The Jacobian also showed up when we *linearized* the Euler equations around a point q in the motivation of the previous lecture.
- In applications (e.g. root finding), the Jacobian needs to be *evaluated* at some point of interest x . To do this, we computed an analytic representation of Jacobian *by hand*.
- What if we cannot do that or have other reasons of not doing it?

THE FINITE-DIFFERENCE METHOD

- Suppose we want to avoid the calculation of the analytic form for derivative. For the single-variate scalar function $f(x)$ from the previous example, we found the following relationship through the Taylor series expansion:

$$f(x + \varepsilon) = f(x) + \left. \frac{df}{dx} \right|_x \varepsilon + \text{h.o.t.},$$

where ε is a *small parameter*.

- If we again drop the higher order terms, we get the following *finite difference* approximation for the derivative, where ε is a *characteristic length scale* with units of x :

$$\left. \frac{df}{dx} \right|_x \approx \frac{f(x + \varepsilon) - f(x)}{\varepsilon}$$

THE FINITE-DIFFERENCE METHOD

- We have *introduced another parameter* in order to approximate the derivative *numerically* with knowledge of $f(x)$ only.
- We *do not know how to choose ε* but we know it has to be small because our Taylor series Ansatz assumes we are looking *in the close vicinity of point x* .
- Let us assume a value $\varepsilon = 10^{-2}$ and replace the previous exact Jacobian in the Newton module:

```
1 import numpy as np
2
3 f = lambda x: x - np.exp(-2.0 * np.sin(4.0 * x) * np.sin(4.0 * x))
4 J = lambda x, eps: (f(x + eps) - f(x)) / eps # Finite-Difference approximation of J
```

THE FINITE-DIFFERENCE METHOD

- Let us assume a value $\varepsilon = 10^{-2}$ and replace the previous exact Jacobian in the Newton module:

```
1 import numpy as np
2
3 f = lambda x: x - np.exp(-2.0 * np.sin(4.0 * x)) * np.sin(4.0 * x)
4 J = lambda x, eps: (f(x + eps) - f(x)) / eps # Finite-Difference approximation of J
```

- If we run Newton's method again, with the numerical approximation of $J(x)$: (see lecture code handouts)


```
1 >>> from newton_fd import (f, J, newton)
2 >>> root = newton(f, J, 0.1, eps=1.0e-2)
3 Iteration 1: Delta x = 1.211561e-01
4 Iteration 2: Delta x = 2.482629e-02
5 Iteration 3: Delta x = 1.424802e-03
6 Iteration 4: Delta x = -4.341516e-05
7 Iteration 5: Delta x = 1.539820e-06
8 Iteration 6: Delta x = -5.437925e-08
9 Found root 2.473652e-01 at iteration 7 # compared to before: 2 extra iterations!
10 >>> f(root)
11 1.8454707206849719e-10 # compared to before: 7 orders of magnitude less accurate!
```

THE FINITE-DIFFERENCE METHOD

- Since we know the exact form for $J(x)$ we can analyze the *numerical error* of our Finite-Difference approximation as we vary ε :

```
1 import numpy as np
2
3 f = lambda x: x - np.exp(-2.0 * np.sin(4.0 * x) * np.sin(4.0 * x))
4 J = lambda x: 1.0 + 16.0 * np.exp(-2.0 * np.sin(4.0 * x)**2
5               ) * np.sin(4.0 * x) * np.cos(4.0 * x) # exact J
6 J_fd = lambda x, eps: (
7     f(x + eps) - f(x)
8 ) / eps # Finite-Difference approximation of J
9
10 x = np.linspace(0.0, 2.0, 1000) # domain for f
11 epsilon = np.logspace(-13, 1, 1000) # discretization \epsilon
12 error = [] # array for L2 errors
13 for eps in epsilon:
14     e = J_fd(x, eps) - J(x) # numerical error for all values in x
15     error.append(np.linalg.norm(e)) # compute L2 error norm
```

THE FINITE-DIFFERENCE METHOD



Observations:

- The numerical error for this approximation of $J(x)$ has a minimum at around 10^{-6}
- The minimum error was **not** obtained at the smallest possible ϵ of about 10^{-16} (*machine precision* → for double precision based on to the IEEE 754 standard)
- Too small ϵ amplify the floating point error while ϵ too large does not provide a good approximation for the derivative!
- The method *reduces* the floating point error by one *decade* if we reduce ϵ by one decade (**1st-order accurate**).

THE FINITE-DIFFERENCE METHOD

- It is *not clear* how to choose the best ε in general. Some results from numerical analysis suggest that it should be around $\sqrt{\varepsilon_{\text{machine}}}$ as a *rule of thumb* for a *1st-order method*.
- **Proof:** in the example before, the minimum numerical error was 1.438669×10^{-6} and corresponds to $\varepsilon = 2.860596 \times 10^{-9}$. If we compute the square root of $\varepsilon_{\text{machine}}$ in Python we find:

```
1 >>> np.finfo(float).eps
2 2.220446049250313e-16
3 >>> np.sqrt(_)
4 1.4901161193847656e-08
```

which is about the ε we can find in the plot on the previous slide.

TOWARDS AUTOMATIC DIFFERENTIATION

- In the introduction, we motivated the need for computational techniques to compute derivatives.
- We focused on the Jacobian J , a $n \times m$ matrix with first order partial derivatives of a mapping $f(x) : \mathbb{R}^m \mapsto \mathbb{R}^n$.
- We have discussed the computation of J with symbolic math which is accurate but may not always be applicable depending on $f(x)$ or may be too costly to evaluate.
- Numerical computation of J may be an alternative method at the cost of accuracy reduction and possible stability issues.
- Automatic differentiation (AD) overcomes both of these deficiencies. It is less costly than symbolic differentiation while evaluating derivatives at *machine precision*. There are two modes of AD: **forward mode** and **reverse mode**, both involve the Jacobian J . *The back-propagation algorithm in machine learning is a special case of reverse mode AD.*

THE BASIC IDEAS OF AUTOMATIC DIFFERENTIATION

- We have discussed the computation of the Jacobian J using symbolic math tools which is accurate but may not always be applicable depending on the complexity of $f(x)$ (or it may be too costly to evaluate).
- Numerical computation of J may be an alternative method at the cost of accuracy and possible stability issues.
- **Automatic differentiation (AD)** overcomes both of these deficiencies. It is:
 - less costly than symbolic differentiation
 - evaluates derivatives to machine precision
- There are two modes of AD: **forward mode** and **reverse mode**. *The back-propagation algorithm in machine learning is a special case of reverse mode automatic differentiation.*
- **Automatic differentiation is based on evaluating the **chain rule** step by step.**

REVIEW OF THE CHAIN RULE

At the heart of AD is the piecewise evaluation of the *chain rule*.

- Suppose we have a function $f(u(t))$ and we want to compute the derivative of f with respect to t . This derivative is given by

$$\frac{df}{dt} = \frac{\partial f}{\partial u} \frac{du}{dt}$$

- Example:** given $f(u(t)) = \sin(4t)$ with $u(t) = 4t$:

$$\frac{\partial f}{\partial u} = \cos(u), \quad \frac{du}{dt} = 4 \quad \Rightarrow \quad \frac{df}{dt} = 4 \cos(4t)$$

REVIEW OF THE CHAIN RULE

The **total change** of f is given by the sum of the partial changes in each coordinate direction.

- Suppose f has another coordinate $v(t)$ so that we have $f(u(t), v(t))$. Once again, we want to compute the derivative of f with respect to t . Applying the chain rule in this case gives

$$\frac{df}{dt} = \underbrace{\frac{\partial f}{\partial u} \frac{du}{dt}}_{\text{Change due to } u} + \underbrace{\frac{\partial f}{\partial v} \frac{dv}{dt}}_{\text{Change due to } v}$$

- Later we will extend this to an **arbitrary** number of coordinates u, v, w, \dots

REVIEW OF THE CHAIN RULE

- *Examples:*

$$\frac{df}{dt} = \underbrace{\frac{\partial f}{\partial u} \frac{du}{dt}}_{\text{Change due to } u} + \underbrace{\frac{\partial f}{\partial v} \frac{dv}{dt}}_{\text{Change due to } v}$$

$$f(u(t), v(t)) = u + v \quad \Rightarrow \quad \frac{df}{dt} = \frac{du}{dt} + \frac{dv}{dt}$$

$$f(u(t), v(t)) = uv \quad \Rightarrow \quad \frac{df}{dt} = v \frac{du}{dt} + u \frac{dv}{dt}$$

$$f(u(t), v(t)) = \sin(uv) \quad \Rightarrow \quad \frac{df}{dt} = v \cos(uv) \frac{du}{dt} + u \cos(uv) \frac{dv}{dt}$$

REVIEW OF THE CHAIN RULE

The gradient operator ∇ :

In vector calculus, the gradient describes the *steepest increase* of a scalar function $f(x)$. This steepest increase is along a certain *direction* given by coordinates $x \in \mathbb{R}^m$.

In our 3D world the dimension is $m = 3$. In general, the coordinate x is can be m -dimensional. Assuming a 3D space with coordinates $x = [x_1, x_2, x_3]^\top$, the *gradient operator* is given by

$$\nabla = \left[\frac{\partial}{\partial x_1}, \frac{\partial}{\partial x_2}, \frac{\partial}{\partial x_3} \right]^\top$$

REVIEW OF THE CHAIN RULE

The gradient operator ∇ (back to chain rule):

- What happens if we replace the scalar parameter $t \in \mathbb{R}$ from before with new coordinates $x \in \mathbb{R}^m$ (a vector of coordinates)?
- Now we compute the *gradient* of a *scalar function* f with respect to all coordinates x and write $f(u(x), v(x))$. To do so we *replace the d/dt operator* from before with the *gradient ∇ operator*:

$$\nabla_x f = \frac{\partial f}{\partial u} \nabla u + \frac{\partial f}{\partial v} \nabla v,$$

where ∇_x on the left side emphasizes that the gradient is *with respect to x* .

- For $u = u(x)$ and $v = v(x)$ it is clear that the only possible gradient is with respect to x .

REVIEW OF THE CHAIN RULE

The gradient operator ∇ (back to chain rule):

Single independent coordinate t (f is a **scalar function**)

$$\frac{df}{dt} = \frac{\partial f}{\partial u} \frac{du}{dt} + \frac{\partial f}{\partial v} \frac{dv}{dt}$$

Vector of independent coordinates x (f is a **scalar function**)

$$\nabla_x f = \frac{\partial f}{\partial u} \nabla u + \frac{\partial f}{\partial v} \nabla v$$

The chain rule still holds, all we did is replace the single coordinate t with an m -dimensional vector of coordinates x . This required us to replace the **differential operator** d/dt with the **differential vector operator** ∇ .

REVIEW OF THE CHAIN RULE

The gradient operator ∇ (back to chain rule):

$$\nabla_x f = \frac{\partial f}{\partial u} \nabla u + \frac{\partial f}{\partial v} \nabla v$$

Example:

- Let $x = [x_1, x_2]^\top \in \mathbb{R}^2$ (two independent coordinates):
 - $u = u(x) = x_1 x_2$
 - $v = v(x) = x_1 + x_2$
- The function is given by: $f(u, v) = \sin(u) - \cos(v)$

$$\nabla u = \begin{bmatrix} x_2 \\ x_1 \end{bmatrix}, \nabla v = \begin{bmatrix} 1 \\ 1 \end{bmatrix} \Rightarrow \nabla_x f = \cos(x_1 x_2) \begin{bmatrix} x_2 \\ x_1 \end{bmatrix} + \sin(x_1 + x_2) \begin{bmatrix} 1 \\ 1 \end{bmatrix}$$

REVIEW OF THE CHAIN RULE

The (almost) general chain rule:

- We now further generalize to not only account for $u = u(x)$ and $v = v(x)$ but *many functions* $y(x) = [y_1(x), \dots, y_n(x)]^\top$ where all y_i take arguments $x \in \mathbb{R}^m$.
- The function $f = f(y(x))$ is a *scalar function* (therefore "almost" general chain rule) of n other functions y_i , each themselves a function of m variables. The gradient of f can then be written as the chain rule of n partial terms:

$$\nabla_x f = \sum_{i=1}^n \frac{\partial f}{\partial y_i} \nabla y_i(x)$$

- *Example on previous slide:* $m = 2$ and $n = 2 \rightarrow y_1 = u = x_1 x_2$ and $y_2 = v = x_1 + x_2$.

EVALUATION (FORWARD) TRACE OF A FUNCTION

- After the chain rule discussion above, let us apply the notation introduced and look at the evaluation trace of a scalar function $f(x)$ with a single argument $x \in \mathbb{R}$ ($m = 1$).

- Consider again the same function from the previous lecture:

$$f(x) = x - \exp(-2(\sin(4x))^2).$$

- We would like to **evaluate the function** at an arbitrary point x_1 . Let us define this point to be $x_1 = \frac{\pi}{16}$.

EVALUATION (FORWARD) TRACE OF A FUNCTION

The evaluation of $f(x_1)$ involves a *partial ordering* of the operations associated with the function f .

$$f(x) = x - \exp(-2(\sin(4x))^2)$$

- **For example:** before we can evaluate $\sin(4x)$ we must evaluate the *intermediate* result $4x$ and before we can evaluate the exponential function we must evaluate the intermediate result $-2(\sin(4x))^2$.
- The evaluation trace introduces *intermediate results* v_j for $j = 1, 2, \dots$ of elementary binary operations like multiplying two numbers together or unary operations like computing $\sin(v_j)$.
- \rightarrow we will use the notation v_1 for the first intermediate result, v_2 for the second and so on.

EVALUATION (FORWARD) TRACE OF A FUNCTION

A word on notation: the coordinates $x = [x_1, \dots, x_m]^\top$ that is $x \in \mathbb{R}^m$ are called **independent** variables, whereas the **intermediate results** v_j are **dependent** variables, they depend on x .

We further define the independent variables as $v_{k-m} = x_k$ for $k = 1, 2, \dots, m$ in the following evaluation trace.

For example: if $x \in \mathbb{R}^2$ then we would use the notation $v_{-1} = x_1$ and $v_0 = x_2$, whereas v_1 is the first intermediate result of the evaluation trace, v_2 the second and so on.

Recall: $f(x_1) = x_1 - \exp(-2(\sin(4x_1))^2)$ and we are interested in the value of $f(x_1 = \frac{\pi}{16})$:

EVALUATION (FORWARD) TRACE OF A FUNCTION

Recall: $f(x_1) = x_1 - \exp(-2(\sin(4x_1))^2)$ and we are interested in the value of $f(x_1 = \frac{\pi}{16})$:


Intermediate	Elementary Operation	Numerical value
$v_0 = x_1$	$\frac{\pi}{16}$	1.963495e-01
v_1	$4v_0$	7.853982e-01
v_2	$\sin(v_1)$	7.071068e-01
v_3	v_2^2	5.000000e-01
v_4	$-2v_3$	-1.000000e+00
v_5	$\exp(v_4)$	3.678794e-01
v_6	$-v_5$	-3.678794e-01
$v_7 = f(x_1)$	$v_0 + v_6$	-1.715299e-01

Input variables (*independent* variables)

Intermediate variables (*dependent* variables, $v_j = v_j(x)$)

COMPUTATIONAL (FORWARD) GRAPH

We can think of each intermediate result v_j as a *node in a graph*. By doing so, we can get a visual interpretation of the partial ordering of *elementary operations* in $f(x) = x - \exp(-2(\sin(4x))^2)$:



COMPUTATIONAL (FORWARD) GRAPH

The *first key observation* is that we worked **from the inside out** when developing the forward evaluation trace. We started from the value we want to evaluate $x_1 = \frac{\pi}{16}$ and worked from the inside outwards to the actual function value $f(x_1)$. The *second key observation* is that in each evaluation step, we only carried out **elementary operations between intermediate results v_j** .

(Later when we look at the reverse mode we will observe that this mode works in the opposite direction)

COMPUTING THE DERIVATIVE IN FORWARD TRACE

We are half-way through the forward mode of automatic differentiation:

- We have identified a partial ordering of elementary operations when evaluating an arbitrary function f .
- By breaking down the problem into smaller parts (elementary functions), we have computed intermediate results v_j for $j = 1, 2, \dots$ with each $v_j = v_j(x)$ evaluated at point $x = x_1$.
- We have associated each v_j to a *node in a graph* for a visualization of the partial ordering. (Try to think about that in terms of a *data structures* → for example a *tree*.)
- We did not compute any derivative so far.

COMPUTING THE DERIVATIVE IN FORWARD TRACE

Let us now return to the gradient ∇ :

In the forward mode of automatic differentiation, we evaluate and carry forward a *directional derivative* of each intermediate variable v_j in a given direction $p \in \mathbb{R}^m$, *simultaneously* with the evaluation of v_j itself. (The latter is what we just did above.)

What does "direction" mean:

- Recall the linearization of the Euler equations (lecture 10): the *direction* was the one that gave the *best linear approximation*.
- Heat flows from *hot to cold*: the "direction" is *dictated* by the *second law of thermodynamics* and given by the temperature gradient $\rightarrow -\nabla T$.
- In forward mode AD: the *direction* is the one of a particular *derivative* we are interested in. It can be *any linear combination of derivatives*. The direction here is a *parameter* that is chosen by you (not a physical law).

COMPUTING THE DERIVATIVE IN FORWARD TRACE

Directional derivative:

Let us **define** the gradient operator in a slightly different way than we did before. Instead of working with the *gradient vector*, we **project** it in the **direction of a vector p** (*the seed vector*):

$$D_p y_i \stackrel{\text{def}}{=} (\nabla y_i)^\top p = \sum_{j=1}^m \frac{\partial y_i}{\partial x_j} p_j.$$

Is the quantity $D_p y_i$ a **vector** or a **scalar**?

COMPUTING THE DERIVATIVE IN FORWARD TRACE

$$D_p y_i \stackrel{\text{def}}{=} (\nabla y_i)^\top p = \sum_{j=1}^m \frac{\partial y_i}{\partial x_j} p_j.$$

We now return to our example function from before:

$$f(x) = x - \exp(-2(\sin(4x))^2),$$

evaluated at the point $x = x_1 = \frac{\pi}{16}$. Recall that $x \in \mathbb{R}$ is one-dimensional ($m = 1$):

We can only choose one possible direction in this case!

Most natural choice is $p = 1$ and simplify:

$$D_p y_i = \frac{\partial y_i}{\partial x_1}$$

COMPUTING THE DERIVATIVE IN FORWARD TRACE

In forward mode AD, the derivative we compute is $D_p y_i$.

The y_i in $D_p y_i$ are just functions that depend on x (the independent variables).

Let us **define** these functions as follows:

$$y_i = v_{i-m} \quad \text{for } i = 1, 2, \dots, n,$$

where the v_j are the variables we used in the evaluation trace before. The number n is the sum of **independent** variables and **intermediate** variables v_j for which $j > 0$.

In the forward trace example from earlier, we have $n = 8$ (1 independent variable and 7 intermediate variables). We therefore get $\rightarrow y_1 = v_0, y_2 = v_1$ up to $y_8 = v_7$. **Note that y_i for $i = 1, 2, \dots, m$ are exactly the m independent variables.**

COMPUTING THE DERIVATIVE IN FORWARD TRACE

Before we compute the forward trace including the derivatives of intermediate variables v_j , let us pick a few examples and look at them individually. Here is where you see how the *chain rule* enters forward mode AD.

Note: from now on we no longer write ∇_x to indicate that the differentials are with respect to the independent coordinates x (think of them as input variables).
→ *we assume this is always the case.*

COMPUTING THE DERIVATIVE IN FORWARD TRACE

What is the value of $D_p v_0$ in the evaluation trace of before?

- We know that $v_0 = x_1$.
- We have chosen our direction parameter to be $p = 1$.
- We are dealing with a 1D problem \rightarrow the gradient operator simply is $\nabla := \partial/\partial x_1$.
- Putting this together we find:

$$D_p v_0 = (\nabla v_0)^\top p = \frac{\partial x_1}{\partial x_1} \cdot 1 = 1$$

COMPUTING THE DERIVATIVE IN FORWARD TRACE

What is the value of $D_p v_2$ in the evaluation trace of before?

- We know that $v_2 = v_2(v_1) = \sin(v_1)$.
- We have chosen our direction parameter to be $p = 1$.
- Because all v_j are functions of the *independent coordinates x* , *the chain rule must be applied here:*

$$\nabla v_2 = \frac{\partial v_2}{\partial v_1} \nabla v_1 = \cos(v_1) \nabla v_1$$

- Putting this together we find:

$$D_p v_2 = (\nabla v_2)^\top p = \cos(v_1) (\nabla v_1)^\top p = \cos(v_1) D_p v_1$$

- **Observe:** we can compute the derivative of v_j with knowledge of v_i and $D_p v_i$ for $i < j$!

COMPUTING THE DERIVATIVE IN FORWARD TRACE

What is the value of $D_p v_7$ in the evaluation trace of before?

- We know that $v_7 = v_7(v_0, v_6) = v_0 + v_6$.
- Nothing new here **except** that we further know $v_7 = f(x_1)$ such that $\nabla v_7 = \nabla f$ and the directional derivative $D_p v_7 = D_p f$ is *exactly the derivative we are after, evaluated at x_1* :

$$\nabla v_7 = \frac{\partial v_7}{\partial v_0} \nabla v_0 + \frac{\partial v_7}{\partial v_6} \nabla v_6.$$

- Projection in direction of p yields:

$$D_p f = D_p v_7 = D_p v_0 + D_p v_6$$

(**Note:** the partial derivatives $\partial v_7 / \partial v_0$ and $\partial v_7 / \partial v_6$ are both equal to 1 in this case)

COMPUTING THE DERIVATIVE IN FORWARD TRACE

We now repeat the computation of the forward trace for our test function $f(x)$. What we did earlier is called the *forward primal trace*, we extend it this time with the *forward tangent trace* which corresponds to the *derivatives of the intermediate variables*.

In forward mode of automatic differentiation, we evaluate and *carry forward a directional derivative* $D_p v_j$ of each intermediate variable v_j in a given direction $p \in \mathbb{R}^m$, *simultaneously* with the evaluation of v_j itself.

Recall: $f(x_1) = x_1 - \exp(-2(\sin(4x_1))^2)$ and we are interested in the value of $\left. \frac{\partial f}{\partial x} \right|_{x_1=\pi/16}$:

AUTOMATIC DIFFERENTIATION: FORWARD MODE

Recall: $f(x_1) = x_1 - \exp(-2(\sin(4x_1))^2)$ and we are interested in the value of $\frac{\partial f}{\partial x} \Big|_{x_1=\pi/16}$:

Forward primal trace	Forward tangent trace	Numerical value: $v_j; D_p v_j$
$v_0 = x_1 = \frac{\pi}{16}$	$D_p v_0 = 1$	1.963495e-01; 1.000000e+00
$v_1 = 4v_0$	$D_p v_1 = 4D_p v_0$	7.853982e-01; 4.000000e+00
$v_2 = \sin(v_1)$	$D_p v_2 = \cos(v_1)D_p v_1$	7.071068e-01; 2.828427e+00
$v_3 = v_2^2$	$D_p v_3 = 2v_2 D_p v_2$	5.000000e-01; 4.000000e+00
$v_4 = -2v_3$	$D_p v_4 = -2D_p v_3$	-1.000000e+00; -8.000000e+00
$v_5 = \exp(v_4)$	$D_p v_5 = \exp(v_4)D_p v_4$	3.678794e-01; -2.943036e+00
$v_6 = -v_5$	$D_p v_6 = -D_p v_5$	-3.678794e-01; 2.943036e+00
$v_7 = f(x_1) = v_0 + v_6$	$D_p v_7 = \frac{\partial f}{\partial x} \Big _{x_1=\pi/16} = D_p v_0 + D_p v_6$	-1.715299e-01; 3.943036e+00

Input variables (*independent* variables)

Intermediate variables (*dependent* variables, $v_j = v_j(x)$)

AUTOMATIC DIFFERENTIATION: FORWARD MODE

Recall: $f(x_1) = x_1 - \exp(-2(\sin(4x_1))^2)$ and we are interested in the value of $\left. \frac{\partial f}{\partial x} \right|_{x_1=\pi/16}$:

Forward primal trace	Forward tangent trace	Numerical value: $v_j; D_p v_j$
$v_0 = x_1 = \frac{\pi}{16}$	$D_p v_0 = 1$	1.963495e-01; 1.000000e+00
$v_7 = f(x_1) = v_0 + v_6$	$D_p v_7 = \left. \frac{\partial f}{\partial x} \right _{x_1=\pi/16} = D_p v_0 + D_p v_6$	-1.715299e-01; 3.943036e+00

We have computed the derivative before on paper and with [sympy](#). You are encouraged to check that we indeed compute the correct result.

AUTOMATIC DIFFERENTIATION: FORWARD MODE

That is all there is to forward mode AD!

The key observations are the following:

- We have broken down the evaluation of an arbitrary function $f(x)$ into smaller pieces, each only consists of *elementary* operations like addition, multiplication, division, subtraction, exponentiation, trigonometric functions and so on.
- Forward mode works **from the inside out**.
- We have computed a **primal trace** of intermediate variables v_j and a **tangent trace** of their directional derivatives $D_p v_j$ both *simultaneously* in the same step.
- Since we only work with **elementary functions**, we know their derivatives and computing $D_p v_j$ is trivial!

RECAP

Automatic Differentiation: Forward Mode (basics)

- Example of Newton's method and numerical schemes to approximate derivatives.
- Evaluation trace
- The computational graph
- Computing derivatives of one variable using the forward mode

Further reading:

- P. H.W. Hoffmann, *A Hitchhiker's Guide to Automatic Differentiation*, Springer 2015, [doi:10.1007/s11075-015-0067-6](https://doi.org/10.1007/s11075-015-0067-6) (You can access this paper through the Harvard network or find it in the class repository)
- Griewank, A. and Walther, A., *Evaluating derivatives: principles and techniques of algorithmic differentiation*, SIAM 2008, Vol. 105
- Nocedal, J. and Wright, S., *Numerical Optimization*, Springer 2006, 2nd Edition
- Finite difference method: https://en.wikipedia.org/wiki/Finite_difference_method