

SYSTEMS DEVELOPMENT FOR COMPUTATIONAL SCIENCE

LECTURE 19

Fabian Wermelinger

Harvard University

CS107 / AC207

Thursday, November 3rd 2022

LAST TIME

- Continuation with binary search trees
- Binary tree traversal
- Other common data structures
- Priority queues and heaps

TODAY

Main topics: *Python generators and coroutines*

Details:

- Generators
- Coroutines
- Python internals: objects, bytecode and interpreter

GENERATORS

- In lecture 17 we discussed the *iterator* design pattern. Iterators are used for the following:

Provide a way to access the elements of an aggregate object *sequentially without exposing its underlying representation*.


- Iterators are *fundamental* for processing of data. So far our conception was that the *data* we are iterating over lives in Random Access Memory (RAM) on your computer (e.g. list, tuple or a linked lists).
- *What if the data we need to process is too large to fit into the RAM?* Then we need a way to *lazily* fetch new elements as we call *next()*.

Every *generator is an iterator* and fully implements the iterator interface. An *iterator* retrieves its items from a collection (e.g. a list), while a *generator* can *produce* items upon request (it is "lazy").

→ the terms "*iterator*" and "*generator*" are often used interchangeably. Be aware of the difference above!

GENERATORS

Recall: the iterator interface (Python obtains *iterators* from *iterables*)



- **Iterable:** an object `x` that is *iterable* implements the `__iter__` special method, which returns *a new* iterator for every `iter(x)` call.
- **Iterator:** an iterator implements the standard interface:
 - `__next__`: returns the next available item. Raises `StopIteration` when there are no more items.
 - `__iter__`: returns `self`. (Allows iterators to be used where an *iterable* is expected.)
- **Generator:** has the same interface as an iterator.

GENERATORS

- **Generators** were added in Python 2.2 in 2001 and introduced the new keyword → **yield**.
- Generators are defined in **PEP 255 -- Simple Generators**.
- A generator **function** is essentially the same as a regular function, except that it **yields** (or produces a value) rather than **returning** from a function (*a generator function can return as well, however*).
- A generator "g" **yields a value whenever we call next(g)** on it. It is then **temporarily suspended** until we call next(g) again.
→ generator functions are **lazy** and execute **on demand**.
- A generator function raises **StopIteration** if either a **return** statement or the end of the generator function body is reached.

GENERATOR FUNCTION

- Any Python function that has a `yield` keyword in its body is a *generator function*.
- A generator function results in a *generator object* when called.
- You can think of generator functions as *factories* for generator objects.

Example:

```
1 def gen_107():
2     """Generator function"""
3     yield 1
4     yield 0
5     yield 7
```

- The generator yields the value that follows every `yield` keyword.
- The `next()` built-in advances the generator the same way it does for iterators. *It will advance to the next yield keyword.*

```
1 >>> from inspect import getgeneratorstate
2 >>> gen_107
3 <function gen_107 at 0x7fe79f4e7310>
4 >>> g = gen_107()
5 <generator object gen_107 at 0x7fe79f57d820>
6 >>> getgeneratorstate(g)
7 'GEN_CREATED'
8 >>> next(g)
9 1
10 >>> getgeneratorstate(g)
11 'GEN_SUSPENDED'
12 >>> list(g)
13 [0, 7]
14 >>> getgeneratorstate(g)
15 'GEN_CLOSED'
```

GENERATOR FUNCTION

Example:

```
1 def gen_107():
2     """Generator function"""
3     yield 1
4     yield 0
5     yield 7
```

```
1 >>> for i in gen_107():
2     ...     i
3     ...
4 1
5 0
6 7
```

- The generator yields the value that follows every `yield` keyword.
- The `next()` built-in advances the generator similarly to iterators.

→ a generator implements the standard iterator interface. *A generator can be used interchangeably with iterators!*

Generator states: (see `inspect.getgeneratorstate`)

State	Description
GEN_CREATED	Waiting to start execution
GEN_RUNNING	Currently being executed by the interpreter
GEN_SUSPENDED	Currently suspended at a <code>yield</code> expression
GEN_CLOSED	Execution has completed

GENERATOR FUNCTION

Example:

```
1 def gen_107():
2     """Generator function"""
3     yield 1
4     yield 0
5     return
6     yield 7 # never executed!
```

```
1 >>> for i in gen_107():
2     ...     i
3     ...
4     1
5     0
```

You can use the `return` keyword inside a generator function. Upon return, the generator will raise `StopIteration` and `yield 7` will *never be reached*.

A generator function creates separate instances of *generator objects* → a generator function is a *factory* for generator objects:

```
1 >>> a, b = gen_107(), gen_107()
2 >>> a; b
3 <generator object gen_107 at 0x7fe79f2cf120>
4 <generator object gen_107 at 0x7fe79f2cf190>
5 >>> iter(a); iter(b)
6 <generator object gen_107 at 0x7fe79f2cf120>
7 <generator object gen_107 at 0x7fe79f2cf190>
```


LINKED LIST ITERATOR REVISITED

Iterator implementation for linked list in lecture 17:

```
1 class LinkedList:
2     # other code skipped
3
4     def __iter__(self):
5         return LinkedListForwardIterator(
6             self.first)
7
8 class LinkedListForwardIterator:
9     def __init__(self, start):
10        self.node = start
11
12    def __next__(self):
13        if self.node != None:
14            curr_node = self.node
15            self.node = self.node.next
16            return curr_node
17        else:
18            raise StopIteration
19
20    def __iter__(self):
21        return self
```

Same functionality implemented with a *generator* function:

```
1 class LinkedList:
2     # other code skipped
3
4     def __iter__(self):
5         node = self.first
6         while node != None:
7             yield node
8             node = node.next
```

- `__iter__` is now a **generator function** (a factory of generators).
- A generator implements the standard iterator interface, otherwise this would not work!
- *No need to write an extra class for it! Less code → way more elegant.*

GENERATOR EXPRESSIONS

Comprehensions:

- Comprehensions provide a concise way to build lists, sets or dictionaries.

- **list comprehension** example:

```
1 >>> [x for x in range(5)]
2 [0, 1, 2, 3, 4]
```

- **set comprehension** example:

```
1 >>> {x for x in [0, 0, 1, 2, 2]}
2 {0, 1, 2}
```

- **dict comprehension** example:

```
1 >>> {k:v for k, v in [(0,'a'), (1,'b'), (2,'c')]}
2 {0: 'a', 1: 'b', 2: 'c'}
```

- Comprehensions are built **eagerly**. Once created, the complete data structure exists in memory (RAM).

Generator expressions:

A generator expression can be thought of as a **lazy** version of a list comprehension.

- It **does not** eagerly build a list, but returns a generator that will **lazily produce the items on demand**.
- A generator expression is written using parentheses **()** instead of brackets **[]** or curly braces **{}**.
- **Generator expression** example:

```
1 >>> (x.upper() for x in list('abc'))
2 <generator object <genexpr> at 0x7f6c2d732430>
```

→ you cannot use the **yield** or **yield from** keywords in a generator expression!

GENERATOR EXPRESSIONS

Iterable via *list comprehension*:

```
1 # generator function
2 def gen():
3     print('Start')
4     yield 'A'
5     print('Continue')
6     yield 'B'
7     print('End')
8
9 # list comprehension (eager)
10 lc = [x for x in gen()]
11
12 # iterate over list
13 print('Enter loop')
14 for i in lc:
15     print(i)
```

Output (eager):

```
1 Start
2 Continue
3 End
4 Enter loop
5 A
6 B
```

Iterable via *generator expression*:

```
1 # generator function
2 def gen():
3     print('Start')
4     yield 'A'
5     print('Continue')
6     yield 'B'
7     print('End')
8
9 # generator expression (lazy)
10 ge = (x for x in gen())
11
12 # iterate over generator expression
13 print('Enter loop')
14 for i in ge:
15     print(i)
```

Output (lazy):

```
1 Enter loop
2 Start
3 A
4 Continue
5 B
6 End
```

GENERATOR EXPRESSIONS

- Generator expressions are syntactic sugar in Python.
- Proposal for generator expressions → [PEP 289](#)
- They are nice in places where you want to be *brief and concise* (same with comprehensions).
- *Generator functions* and *generator expressions* both return *generator objects* → *they are both just generator factories*.
- *Generator functions* are much more flexible! They allow for *multiple statements* and more complex code.
- *Generator functions* can further be used as *coroutines*, a generalization that will be introduced later.

GENERATOR FUNCTIONS IN THE STANDARD LIBRARY

- The Python standard library has many *generator utilities implemented* → <https://docs.python.org/3/library/index.html>
- You should be aware of some commonly used tools in order *not to reinvent the wheel*.
- Useful categories for data transformations include:
 - Filters
 - Maps
 - Merge of inputs
 - Expansion of input into multiple outputs
 - Rearrangements
 - Reductions
- Most of the tools above are available in the `itertools` module. Others are built-in without need of importing modules.

GENERATOR FUNCTIONS IN THE STANDARD LIBRARY

Example: `os.walk`

Example directory structure:

```
1 .
2 └─ file_top
3 └─ oswalk.py
4 └─ subdir1
5     └─ file_subdir1
6         └─ subsubdir1
7             └─ file_subsubdir1
8 └─ subdir2
9     └─ file_subdir2
```

`os.walk` generator in Python:

```
1 >>> import os
2 >>> for path, dirs, files in os.walk('./'):
3 ...     print(f'{path}\n\tdirs: {dirs}\n\tfiles: {files}')
4 ...
5 ./
6     dirs: ['subdir1', 'subdir2']
7     files: ['oswalk.py', 'file_top']
8 ./subdir1
9     dirs: ['subsubdir1']
10    files: ['file_subdir1']
11 ./subdir1/subsubdir1
12    dirs: []
13    files: ['file_subsubdir1']
14 ./subdir2
15    dirs: []
16    files: ['file_subdir2']
```

- `os.walk` is very powerful for exploring data in file systems.
- `os.walk` returns a generator object that yields the current directory `path` along with lists of directories `dirs` and `files` contained within it.

GENERATOR FUNCTIONS IN THE STANDARD LIBRARY

Example: `filter`

- Filter elements in an iterable for which a *predicate function* returns true:

```
filter(predicate_function, iterable)
```

- The return object of a call to `filter` is an iterator.
- Example: filter vowels*

```
1 >>> list(filter(lambda x: x.lower() in 'aeiou', 'What you seek is seeking you.'))
2 ['a', 'o', 'u', 'e', 'e', 'i', 'e', 'e', 'i', 'o', 'u']
```

or types:

```
1 >>> list(filter(lambda x: isinstance(x, int), [0, 0x1, 0o2, 3.0]))
2 [0, 1, 2]
```

- The inverse of `filter` is provided by `itertools.filterfalse`:

```
1 >>> from itertools import filterfalse
2 >>> list(filterfalse(lambda x: isinstance(x, int), [0, 0x1, 0o2, 3.0]))
3 [3.0]
```

GENERATOR FUNCTIONS IN THE STANDARD LIBRARY

Example: `map`

- Applies a function to every item of an iterable and *yields* the result.
- You can pass n iterables. In this case the function must take n arguments.
- See `itertools.starmap` for an alternative version to unpack function arguments instead.
- *Example:*

```
1 >>> list(map(lambda a, b: (a, b), range(11), list('ABC')))  
2 [(0, 'A'), (1, 'B'), (2, 'C')] # list('ABC') is the shorter iterable!
```

→ note that `map` stops when the shortest input iterator is exhausted!

- There are many more useful generator functions in `itertools`. Be sure to check them out.

APPLICATION EXAMPLE: LAZY READ OF LARGE DATA

- Generators are useful to create *data readers* for working with data sets.
- These readers *expose the iterator pattern*.
- If your input data set is *very large* it may not fit into RAM entirely. In that case you must produce the data *on the fly* → *lazy read* the data from the disk or tape. *A generator is the right tool for this.*
- **Example:** assume you are working with the following function to process data. The input data may either come from an ASCII text file or a NumPy *binary file* (e.g. due to different data collection procedures):

```
1 from itertools import chain
2 def process(*data_iterables): # can pass multiple input generators...
3     val = 0
4     item_count = 0
5     for item in chain.from_iterable(data_iterables): # ...and chain them together
6         val += item # in this example we just sum up the input data values
7         item_count += 1
8         if item_count % 10000 == 0: # print progress
9             print(f'{item_count} items processed from input')
10    return val
```

APPLICATION EXAMPLE: LAZY READ OF LARGE DATA

- *Example:* data processing function

```
1 from itertools import chain
2 def process(*data_iterables): # can pass multiple input generators...
3     val = 0
4     item_count = 0
5     for item in chain.from_iterable(data_iterables): # ...and chain them together
6         val += item # in this example we just sum up the input data values
7         item_count += 1
8         if item_count % 10000 == 0: # print progress
9             print(f'{item_count} items processed from input')
10    return val
```

- The function takes a number of *data generators* (or just iterables) and chains them together to iterate over all data items combined.
Examples for data items: movie frames, pressure fields from simulations, audio samples, sentences, *anything really*.
- Assume the data generators above yield just one number at a time.
- What if the data file(s) are too large for your physical RAM or if you read from a continuous stream? You must use a *generator* in this case.

APPLICATION EXAMPLE: LAZY READ OF LARGE DATA

- Eager reader for binary data specific to some application (*all data is loaded into RAM*):

```
1 def read_binary_eager(fname): # returns iterable
2     with open(fname, 'r') as f: # open file read-only
3         return np.fromfile(f, dtype=float)
```

- Lazy reader for binary data specific to some application:

```
1 def read_binary_lazy(fname): # returns generator object
2     with open(fname, 'r') as f: # open file read-only
3         while item := np.fromfile(f, dtype=float, count=1):
4             yield item
```

(Note: ":"= is called an *assignment expression* and only works for Python 3.8 and later. See <https://peps.python.org/pep-0572/>)

- When you use this code in practice you notice that the lazy reader (generator) is **100x slower**! *Do you have an idea what could be the problem?*

APPLICATION EXAMPLE: LAZY READ OF LARGE DATA

- Lazy reader for binary data specific to some application:

```
1 def read_binary_lazy(fname): # returns generator object
2     with open(fname, 'r') as f: # open file read-only
3         while item := np.fromfile(f, dtype=float, count=1):
4             yield item
```

- *Bandwidth* on a computer is very precious and it is *optimized* for reading more than *just a few bytes at a time!* → in the example reader above *only 8 bytes are read at a time* (one 64-bit float).

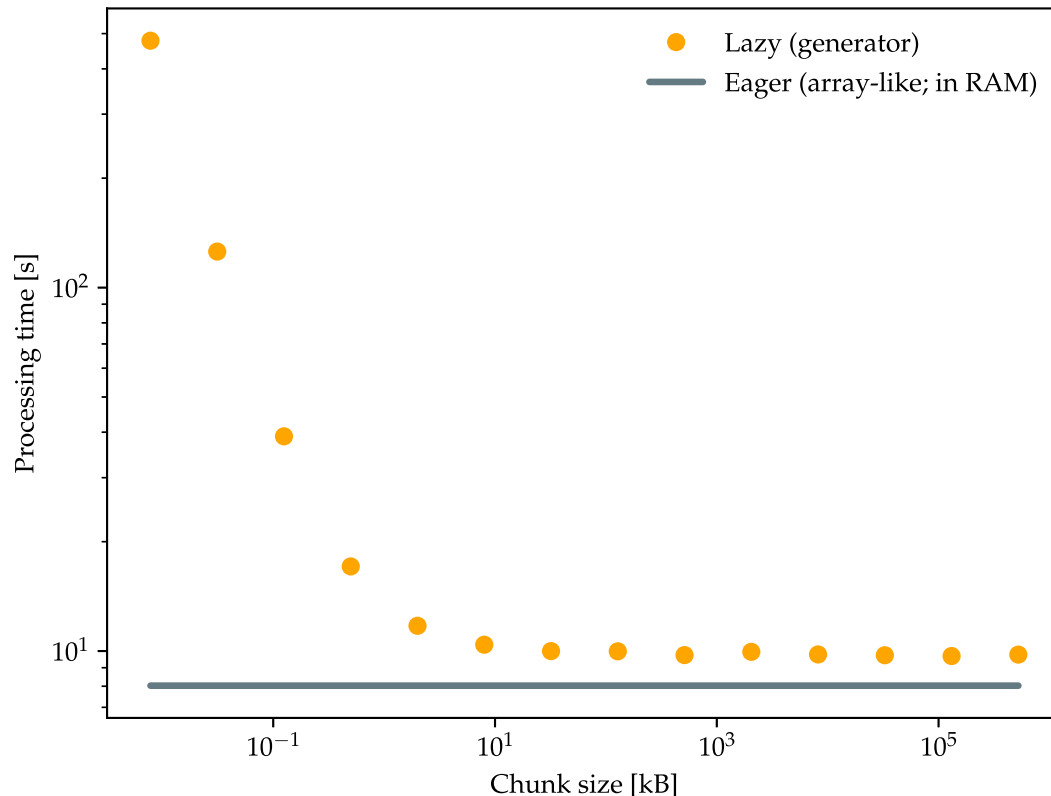
- *Read at least a few kilo-bytes at a time:*

```
1 def read_binary_lazy(fname, chunk_size=1): # returns generator object
2     with open(fname, 'r') as f: # open file read-only
3         while (chunk := np.fromfile(f, dtype=float, count=chunk_size)).size > 0:
4             for item in chunk:
5                 yield item
```

→ *this will saturate the memory bandwidth and improve performance!*

APPLICATION EXAMPLE: LAZY READ OF LARGE DATA

Benchmark for lazy load of 512MB file with different chunk size:



- If you read very small data lazily, you will *not saturate the memory bandwidth* and process up to 100x slower!
- You should read *at least 10–100 kilobyte* at a time. This may depend on the architecture you are running on. The larger the chunk size the better.
- Prefer to eager load the full data set if you can afford it.

COROUTINES

- You are already familiar with the concept of a *function*. They are also called *subroutines* (especially in Fortran):
 - Functions allow you to avoid *code duplication*.
 - They form a logical segmentation of the problem. Also enable easier debugging.
 - When you call a function, temporary variables are allocated (called automatic variables) that exist during the lifetime of the function only. Examples are function arguments or local variables in the function body.
 - A function has one entry and may have multiple return points. It is *asymmetric*. Once you return, the memory for the temporary variables is released.
- A function is a special case of a more general concept called a *coroutine*.
 - Functions are *asymmetric* between caller and callee.
 - Coroutines are *symmetric* between caller and callee → *you can enter and leave a coroutine many times*.

You may think of two coroutines as a "*team*" of programs that repetitively call each other with different input each time.

COROUTINES

- *A coroutine is syntactically the same as a generator:* a function with a **yield** keyword in its body.
- In a coroutine the **yield** keyword usually appears *on the right side of an expression. Example:*

```
1 def coroutine():
2     while True:
3         x = yield # yield may obtain something and we assign it to x!
4         print(x)
```

- **Note:** **yield may or may not** produce a value! If there is no expression after **yield** → it will yield **None** (like in the example code above).
- Unlike a generator, a coroutine **can receive data from the caller** by calling **c.send(data)** (**c.send(data)** returns what **next(c)** would return):

```
1 >>> c = coroutine()
2 >>> next(c) # we must prime the coroutine by calling next() before use
3 >>> y = c.send('Hello CS107/AC207')
4 Hello CS107/AC207
5 >>> print(y)
6 None
```

COROUTINE EXAMPLE

```
1 from inspect import getgeneratorstate
2
3 def coroutine():
4     ncalls = 0
5     while True:
6         x = yield ncalls
7         ncalls += 1
8         print(f'coroutine(): {x} (call id: {ncalls})')
9
10 def main():
11     c = coroutine()
12     print(getgeneratorstate(c))
13     next(c) # prime the coroutine; next() returns 0 here
14     print(getgeneratorstate(c))
15
16     c.send('CS107')
17     print('main(): control back in main function')
18     last_call = c.send('AC207')
19     print(f'main(): called coroutine() {last_call} times')
20
21     c.close()
22     print(getgeneratorstate(c))
23
24 if __name__ == "__main__":
25     main()
```

```
1 GEN_CREATED
2 GEN_SUSPENDED
3 coroutine(): CS107 (call id: 1)
4 main(): control back in main function
5 coroutine(): AC207 (call id: 2)
6 main(): called coroutine() 2 times
7 GEN_CLOSED
```


- line 11 creates the coroutine (nothing has been run yet).
- line 13 primes the coroutine. This means activate it and run until the first yield, then suspend it.
- line 16 sends data to the suspended coroutine. This will activate it and run until the next yield is reached. The `.send()` call is similar to `next()` except that we also send data.
- Coroutines can be shutdown using the `.close()` method in case it never reaches a return statement (or end of function).

GENERATORS/COROUTINES RECAP

- A generator can pause at a **yield** statement. It yields a value back to the caller. The generator state is suspended until `next()` is called on it again.
→ think of `yield` as control flow.
- A generator implements the standard iterator interface.
- Coroutines are an advanced programming concept that involves entering and returning from functions which are in an *intermediate* execution state. A Python coroutine **extends** a generator with `.send`, `.close` and `.throw` methods. **Coroutines and generators are conceptually very different.**
- **Guido van Rossum**: there are three different styles of code you can write using generators:
 1. The traditional "**pull**" style (**example**: linked list `__iter__` we discussed, see [PEP 255](#) and [PEP 380](#)).
 2. The "**push**" style (i.e., the coroutines that we discussed here using `.send` to push, see [PEP 342](#)).
 3. Concurrent tasks (coroutines with `async` and `await` syntax for concurrent programming, see [PEP 492](#)). We do not discuss concurrent programming in this class.

PYTHON INTERNALS: OBJECTS AND FRAMES

Recall: the sketches we saw during the [pythontutor](#) examples



- **Frames:** "frame objects" that execute code (imagine a *stack* data structure: the blue shaded frame is at the top of the stack).
- **Objects:** any other Python objects. Functions, classes, data structures, etc.

- Frame objects *execute* code and form a *sequence* in a *stack* data structure.
- Arrows indicate *references to objects in memory*.
- When we enter a function `f()`, a new frame is *pushed* onto the stack and executes.
- When done, the function frame is *popped* off the stack and we return to the caller frame (global frame here).

- The data structure used to organize frames is a LIFO stack. **Will that work for coroutines?**

PYTHON INTERNALS: OBJECTS

*All the data stored in a Python program is built around the concept of an **object**.*

Terminology:

- Every piece of data is stored in an **object**. This includes Python frames and code.
- Each object has an **identity**, a **type** (also known as its class) and a **value**.
- The identity of an object is its location in memory. **Names** store a **reference** to that a specific memory location.
- The **type** of an object describes the internal representation as well as methods and operations it supports → implemented in a `class`.
- When an object of a specific type is created, we called it an **instance** of that type. *After an instance is created, its identity and type can no longer be changed.*
- If an object's value can be modified, we call it **mutable**, otherwise it is said to be **immutable**.
- **Containers** or **collections** are objects that contain references to other objects.
- Because everything in Python is represented by objects, they are said to be **first class**.

PYTHON INTERNALS: OBJECTS

Example: user defined function object

- User defined functions are *callable* objects created at the module level by using *def* or *lambda*. Functions are *first class* objects in Python.

```
1 >>> def f():
2 ...     pass
3 ...
4 >>> g = lambda x: x
5 >>> f.__code__; g.__code__
6 <code object f at 0x7fd88a3fac90, file "<stdin>", line 1>
7 <code object <lambda> at 0x7fd88a3fabe0, file "<stdin>", line 1>
```

- A user-defined function *f* has the following attributes:

Attribute	Description
<code>f.__doc__</code>	Documentation string
<code>f.__name__</code>	Function name
<code>f.__dict__</code>	Dictionary containing function attributes
<code>f.__code__</code>	Byte-compiled code
<code>f.__defaults__</code>	Tuple containing the default arguments
<code>f.__globals__</code>	Dictionary defining the global namespace
<code>f.__closure__</code>	Tuple containing data related to nested scopes

- Python code *is compiled* into *bytecode* objects on the fly.
- Python is an *interpreted* language. Under the hood, however, code is transformed into bytecode objects. *The interpreter is a virtual machine.*
- Running your code for the first time is slower due to bytecode generation. The result is *cached* in *.pyc* files for faster subsequent execution.

RECAP

- Generators
- Coroutines
- Python internals: objects, bytecode and interpreter

Further reading:

- Generator and yield expressions:
<https://docs.python.org/3/reference/expressions.html#generator-expressions>
- Chapters 14 and 16 in Luciano Ramalho, "*Fluent Python: Clear, Concise, and Effective Programming*", O'Reilly Media, 2015