

SYSTEMS DEVELOPMENT FOR COMPUTATIONAL SCIENCE

LECTURE 20

Fabian Wermelinger

Harvard University

CS107 / AC207

Tuesday, November 8th 2022

LAST TIME

- Generators
- Coroutines

TODAY

Main topics: *Python internals and the interpreter loop*

Details:


- Python internals:
 - Code objects and bytecode
 - The interpreter and the evaluation loop
 - Frame objects
 - Generator objects
- Why are Python built-in lists slow and NumPy arrays fast?

AGENDA CHECK:

- **Quiz 3** takes place on Thursday. *You will have 25 minutes* for 12 questions. The quiz is available within a 12 hour time window (starting 9:00am). Question topics: <https://edstem.org/us/courses/24296/discussion/2039027>
- Milestone 2B due on Thursday. The milestone consists of a brief progress report <https://harvard-iacs.github.io/2022-CS107/project/M2B/>. Please limit to no more than 1/2 page.

PYTHON INTERNALS: OBJECTS AND FRAMES

Recall: the sketches we saw during the [pythontutor](#) examples



- **Frames:** "frame objects" that execute code (imagine a *stack* data structure: the blue shaded frame is at the top of the stack).
- **Objects:** any other Python objects. Functions, classes, data structures, etc.

- Frame objects *execute* code and form a *sequence* in a *stack* data structure.
- Arrows indicate *references to objects in memory*.
- When we enter a function `f()`, a new frame is *pushed* onto the stack and executes (blue shaded frame on left).
- When done, the function frame is *popped* off the stack and we return to the caller frame (global frame on left).

- The data structure used to organize frames is a LIFO stack. ***Will that work for coroutines?***

PYTHON INTERNALS: OBJECTS

*All the data stored in a Python program is built around the concept of an **object**.*

Terminology:

- Every piece of data is stored in an **object**. This includes Python frames and code.
- Each object has an **identity**, a **type** (also known as its class) and a **value**.
- The identity of an object is its location in memory. **Names** store a **reference** to a specific memory location.
- The **type** of an object describes the internal representation as well as methods and operations it supports → implemented in a `class`.
- When an object of a specific type is created, we called it an **instance** of that type. *After an instance is created, its identity and type can no longer be changed.*
- If an object's value can be modified, we call it **mutable**, otherwise it is said to be **immutable**.
- **Containers** or **collections** are objects that contain references to other objects.
- Because everything in Python is represented by objects, they are said to be **first class**.

PYTHON INTERNALS: OBJECTS

Example: user defined function object

- User defined functions are *callable* objects created at the module level by using *def* or *lambda*. Functions are *first class* objects in Python.

```
1 >>> def f():
2 ...     pass
3 ...
4 >>> g = lambda x: x
5 >>> f.__code__; g.__code__
6 <code object f at 0x7fd88a3fac90, file "<stdin>", line 1>
7 <code object <lambda> at 0x7fd88a3fabe0, file "<stdin>", line 1>
```

- A user-defined function *f* has the following attributes:

Attribute	Description
<code>f.__doc__</code>	Documentation string
<code>f.__name__</code>	Function name
<code>f.__dict__</code>	Dictionary containing function attributes
<code>f.__code__</code>	Byte-compiled code
<code>f.__defaults__</code>	Tuple containing the default arguments
<code>f.__globals__</code>	Dictionary defining the global namespace
<code>f.__closure__</code>	Tuple containing data related to nested scopes

- Python code *is compiled* into *bytecode* objects on the fly.
- Python is an *interpreted* language → under the hood, however, code is transformed into *bytecode objects*. The interpreter is a virtual machine.
- Running your code for the first time is slower due to bytecode generation. The result is *cached* in *.pyc* files for faster subsequent execution.

PYTHON INTERNALS: CODE OBJECTS

- *The Python interpreter executes code objects. They represent raw bytecode.*
- We can generate code objects with the `compile()` built-in function:

```
1 >>> a = 1
2 >>> co = compile('a + 1', '<string>', mode='eval')
3 >>> eval(co) # evaluate the code object
4 2
```

- The *raw bytecode* is contained in `co_code`:

```
1 >>> co.co_code
2 b'e\x00d\x00\x17\x00S\x00' # raw binary encoded Python bytecode
```

- We can *disassemble* bytecode into the *instructions* that Python executes for a particular code object:

```
1 >>> import dis
2 >>> dis.dis(co)
3      1           0 LOAD_NAME           0 (a)
4      2           2 LOAD_CONST        0 (1)
5      4           4 BINARY_ADD
6      6           6 RETURN_VALUE
```

4 instructions are executed: 2 loads, 1 binary addition and returning the result. A list of all bytecode instructions can be found [here](#).

PYTHON INTERNALS: INTERPRETER

- All the data stored in a Python program is built around the concept of an *object*. Code objects are *compiled bytecode*. The interpreter turns those code objects into *frame objects* and executes them (*left column in [pythontutor](#)*).
- Bytecode is an implementation detail of the CPython interpreter and **not portable** between different Python versions → *ignore `__pycache__` and `.pyc` files in your Git repositories!*
- → to execute frame objects, *input* data is required (*right column in [pythontutor](#)*).
- The CPython interpreter obtains input data from a *value stack* and executes frame objects arranged in a *frame stack* in a central loop called the *evaluation loop*. In the *interactive* Python shell this is called **REPL**: Read, Evaluate, Print, Loop. The CPython interpreter is written in C → <https://github.com/python/cpython>
- At the **very core** of the evaluation loop is the `_PyEval_EvalFrameDefault` function. This is the function that brings everything together and makes your code come to life. *Everything that is executed in Python must go through this function*. Recent addition of specialized instructions in Python 3.11 have improved this function. See [PEP 659](#) for more details.

PYTHON INTERNALS: NAME SCOPES

- We saw that a `LOAD_NAME` instruction pushes the object at given index in `co_names` onto the *value stack*. This instruction obtains the object from the *local scope*. The `LOAD_GLOBAL` instruction would be used to load a name from the *global scope*.
- In Python, the *local* and *global* scopes can be inspected with the `locals()` and `globals()` built-ins, respectively:

```
1 >>> def f(x):
2 ...     l0 = x # l0 and x are local names
3 ...     l1 = g # g is a global name
4 ...     print(f'id(x) = {id(x)}')
5 ...     print(f'id(g) = {id(g)}')
6 ...     print(f'Local vars in f(x): {locals()}')
7 ...     print(f'Global vars in f(x): {globals()}')
8 ...
9 >>> g = 42
10 >>> f(g)
11 id(x) = 140268417435152 # local and global names hold same reference!
12 id(g) = 140268417435152 # local and global names hold same reference!
13 Local vars in f(x): {'x': 42, 'l0': 42, 'l1': 42}
14 Global vars in f(x): {'__name__': '__main__', '__doc__': None, other special attributes,
15 '__builtins__': <module 'builtins' (built-in)>, 'f': <function f at 0x7f92c937fd90>, 'g': 42}
```


PYTHON INTERNALS: NONLOCAL SCOPE

- Recall the `nonlocal` statement of Problem 4 in Homework 2.
- Python code objects have special attributes for names in scopes:
 - `co_varnames`: names of *local variables*.
 - `co_cellvars`: names of local variables that are *referenced by nested functions*.
- The `nonlocal` statement in Python is implemented based on these attributes and the corresponding *instructions*:

```
1 def global(x): # not same 'x'!  
2     def c(y): # closure  
3         global x # not same 'x'!  
4         x -= y # x = x - y  
5         return x  
6     return c
```

```
1 def nlocal(x): # same 'x'!  
2     def c(y): # closure  
3         nonlocal x # same 'x'!  
4         x -= y # x = x - y  
5         return x  
6     return closure
```

```
1 def local(x):  
2     def c(y): # closure  
3         x -= y # x = x - y  
4         return x  
5     return c
```

```
1 Cell vars 'global': ()  
2 Local vars 'global': ('x', 'c')  
3 Local vars 'c': ('y',)  
4 4 0 LOAD_GLOBAL      0 (x)  
5 2 LOAD_FAST          0 (y)  
6 4 INPLACE_SUBTRACT  
7 6 STORE_GLOBAL       0 (x)  
8  
9 5 8 LOAD_GLOBAL       0 (x)  
10 10 RETURN_VALUE
```

```
1 Cell vars 'nlocal': ('x',)  
2 Local vars 'nlocal': ('x', 'c')  
3 Local vars 'c': ('y',)  
4 4 0 LOAD_DEREF        0 (x)  
5 2 LOAD_FAST          0 (y)  
6 4 INPLACE_SUBTRACT  
7 6 STORE_DEREF        0 (x)  
8  
9 5 8 LOAD_DEREF        0 (x)  
10 10 RETURN_VALUE
```

```
1 Cell vars 'local': ()  
2 Local vars 'local': ('x', 'c')  
3 Local vars 'c': ('y', 'x')  
4 3 0 LOAD_FAST         1 (x)  
5 2 LOAD_FAST          0 (y)  
6 4 INPLACE_SUBTRACT  
7 6 STORE_FAST         1 (x)  
8  
9 4 8 LOAD_FAST         1 (x)  
10 10 RETURN_VALUE
```

PYTHON INTERNALS: NONLOCAL SCOPE

What is causing the `UnboundLocalError`:

```
1 def local(x):
2     def c(y): # closure
3         x = x - y # same as x -= y
4         return x
5     return c
```

```
1 Cell vars `local`: ()
2 Local vars `local`: ('x', 'c')
3 Local vars `c`: ('y', 'x')
4 3 0 LOAD_FAST          1 (x)
5   2 LOAD_FAST          0 (y)
6   4 INPLACE_SUBTRACT
7   6 STORE_FAST         1 (x)
8
9  4 8 LOAD_FAST          1 (x)
10 10 RETURN_VALUE
```

- Python assumes `x` is contained in the local scope because it "sees" an assignment to `x`.
- Python can use more efficient instructions this way → `LOAD_FAST`
- This instruction will fail however because `x` *is not bound to the local scope!*
- Python raises a `UnboundLocalError` because it cannot load `x` at bytecode offset 0!
- The `nonlocal` statement will put `x` into the `co_cellvars` tuple and use the `LOAD_DEREF` instruction instead.

PYTHON INTERNALS: NONLOCAL SCOPE

Without assignment → only read x:

```
1 def local(x):
2     def c(y): # closure
3         y = x # read-only x!
4         return y
5     return c
```

```
1 Cell vars `local`: ('x',)
2 Local vars `local`: ('x', 'c')
3 Local vars `c`: ('y',)
4 3 0 LOAD_DEREF          0 (x)
5   2 STORE_FAST          0 (y)
6
7 4 4 LOAD_FAST            0 (y)
8   6 RETURN_VALUE
```

- The name x is non-local but can be referenced from the outer scope.
- Python realizes this automatically and issues the correct **LOAD_DEREF** instruction for the read operation.

PYTHON INTERNALS: INTERPRETER

Important terms for the Python interpreter:

- The evaluation loop (or REPL in the interactive Python shell) will take a *code object* and convert it into a series of *frame objects*.
- Frame objects are executed in a so called *frame stack* (what we saw in [pythontutor](#)).
- The interpreter manages referenced variables in a *value stack*.
- The interpreter has at least one thread but *at most one thread can run at a time*. The Python interpreter uses an internal global interpreter lock (called **GIL**) which prevents *race conditions* and ensures thread safety. The GIL imposes a very strong constraint on multi-threaded execution and was subject to many discussions in the past. [A recent post \(10/07/2021\) on the python-dev mailing list](#) proposes a new design to remove the GIL which would mean a major change in the Python interpreter, a possible change that will be reality in the next Python 4 release.
- **Did you know:** for the first time in 20 years, [Python became the worlds most popular programming language](#) this year → *this was for 2021, Python still is #1 in 2022.*

PYTHON INTERNALS: BACK TO OBJECTS

Everything in Python is an object!

Fixed size object base:

```
1 typedef struct _object {
2     _PyObject_HEAD_EXTRA
3     Py_ssize_t ob_refcnt;
4     PyObject *ob_type;
5 } PyObject; // C code in cpython
```

- `_PyObject_HEAD_EXTRA` is a macro that is usually empty.
- `ob_refcnt` is the *reference* count for the object.
- `ob_type` is a pointer to the type object. Recall that Python is *dynamically typed*.

Variable size object base:

```
1 typedef struct {
2     PyObject ob_base;
3     Py_ssize_t ob_size;
4 } PyVarObject; // C code in cpython
```

- `ob_base` is a fixed size object instance.
- `ob_size` is the number of items in the variable part.
- Containers (e.g. `list`) are objects of this type.

No object in Python is a direct instance of `PyObject`. However, every object in Python can be cast to a `PyObject` (if it is variable size it can be cast to `PyVarObject` in addition).

PYTHON INTERNALS: FRAME OBJECTS

A frame object is a PyObject with the following additional properties:

The PyFrameObject:

```
1 struct _frame {
2     PyObject ob_base;
3     struct _frame *f_back;
4     struct _interpreter_frame *f_frame;
5     PyObject *f_trace;
6     int f_lineno;
7     char f_trace_lines;
8     char f_trace_opcodes;
9     char f_own_locals_memory;
10 } PyFrameObject;
```

- `ob_base` is the base instance (as before).
- `f_back` is a pointer to the previous PyFrameObject towards the caller (enables the frame stack).
- `f_frame` is a pointer to the frame data.
- Other fields are used for debugging.

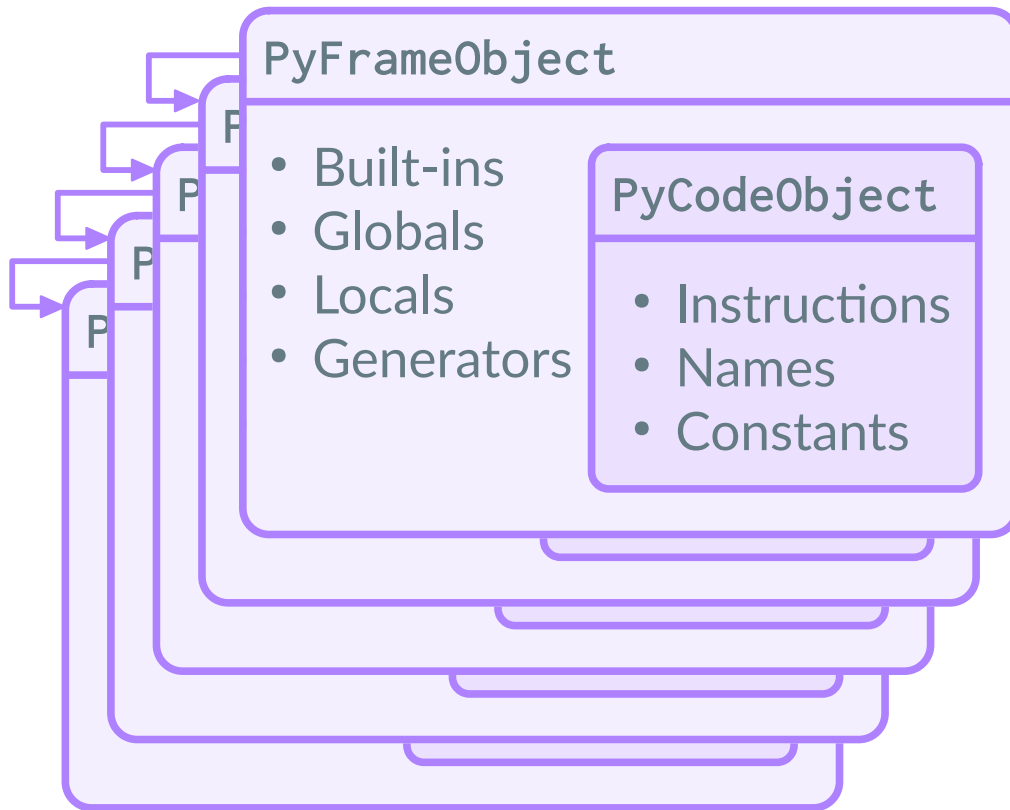
The frame data (some code not shown):

```
1 typedef struct _interpreter_frame {
2     PyObject *f_globals;
3     PyObject *f_builtins;
4     PyObject *f_locals;
5     PyCodeObject *f_code;
6     PyFrameObject *frame_obj;
7     PyObject *generator;
8     int f_lasti;
9     int depth;
10 } InterpreterFrame;
```

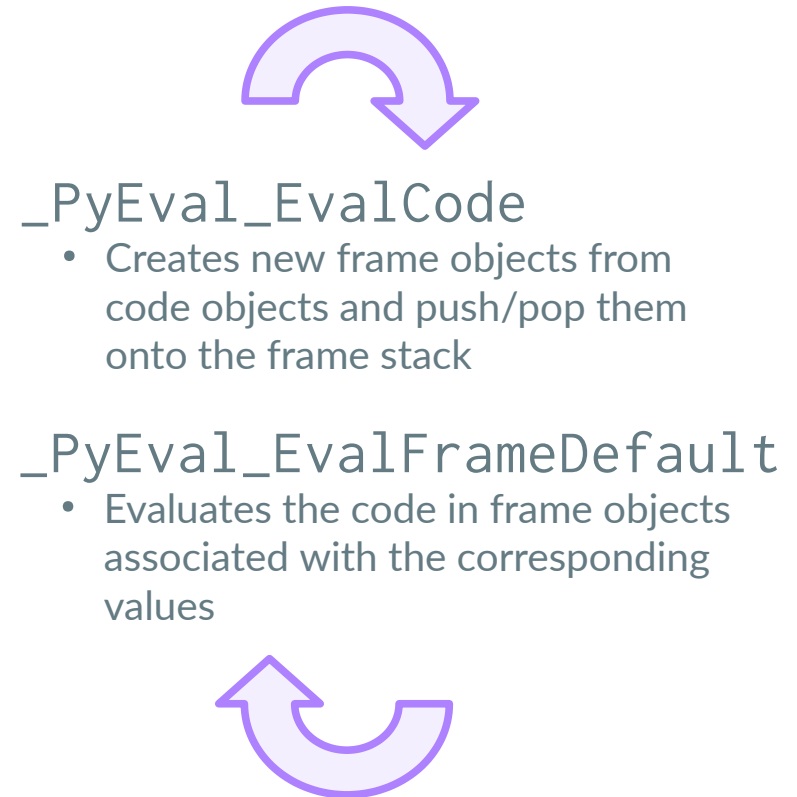
- *Is not an object* (has no `ob_base`!).
- `f_globals` and `f_locals` point to data.
- `f_code` is the bytecode object that will be executed by the frame.
- `f_lasti` index of the last instruction executed → *Where is this index used?*

PYTHON INTERNALS: FRAME OBJECTS

Frame Stack



Evaluation Loop



<https://github.com/python/cpython>

PYTHON INTERNALS: GENERATOR OBJECTS

The `PyGenObject` (some code not shown):

```
1  typedef struct {
2      /* The gi_ prefix is intended to
3       remind of generator-iterator. */
4      PyObject ob_base;
5      /* Note: gi_frame can be NULL if
6       the generator is "finished" */
7      struct _interpreter_frame *gi_xframe;
8      /* The code object backing
9       the generator */
10     PyCodeObject *gi_code;
11 } PyGenObject;
```

- `gi_xframe` points to the current frame object for the generator.
- `gi_code` bytecompiled code object of the *generator function*.
- `PyGenObject`'s are flagged when created with `CO_GENERATOR` (last time), `CO_COROUTINE` (PEP 492) or `CO_ASYNC_GENERATOR` (PEP 525).

- Frame objects have a *pointer to generator objects* and they also store the index of the *last instruction* in the bytecode of the frame.
- The pointer allows to *resume* a generator object that is associated with a frame. The frame data has this code:

```
1  PyObject *generator;
```

This *pointer* points to a `PyGenObject` somewhere in dynamic memory (*not on the frame stack*).

- Although the frame object is in the frame stack, the *generator* pointer allows to obtain the generator object from somewhere else in memory such that it can be resumed (for example when the frame calls `next()` on it).
- **Fact:** Python's frame stack is maintained in *dynamic memory (heap)* → *this does not apply to normal program execution.*

PYTHON INTERNALS: DYNAMIC MEMORY

- Processes share CPU and memory among each other.
- Sharing memory is a non-trivial task when designing operating systems.
- Each physical memory cell (byte-sized cells) can be addressed uniquely. For example:
 - **32-bit system:** 4294967296 addresses; can handle **4GB** (gigabyte) of RAM at most.
 - **64-bit system:** 18446744073709551616 addresses; can handle **16EB** (exabyte) of RAM at most. (This is A LOT!)
- **Virtual memory** simplifies memory management in operating systems by making processes "think" that their memory space always starts at address 0. The **virtual address** is then translated to the real physical address by a hardware component called **memory management unit (MMU)**.

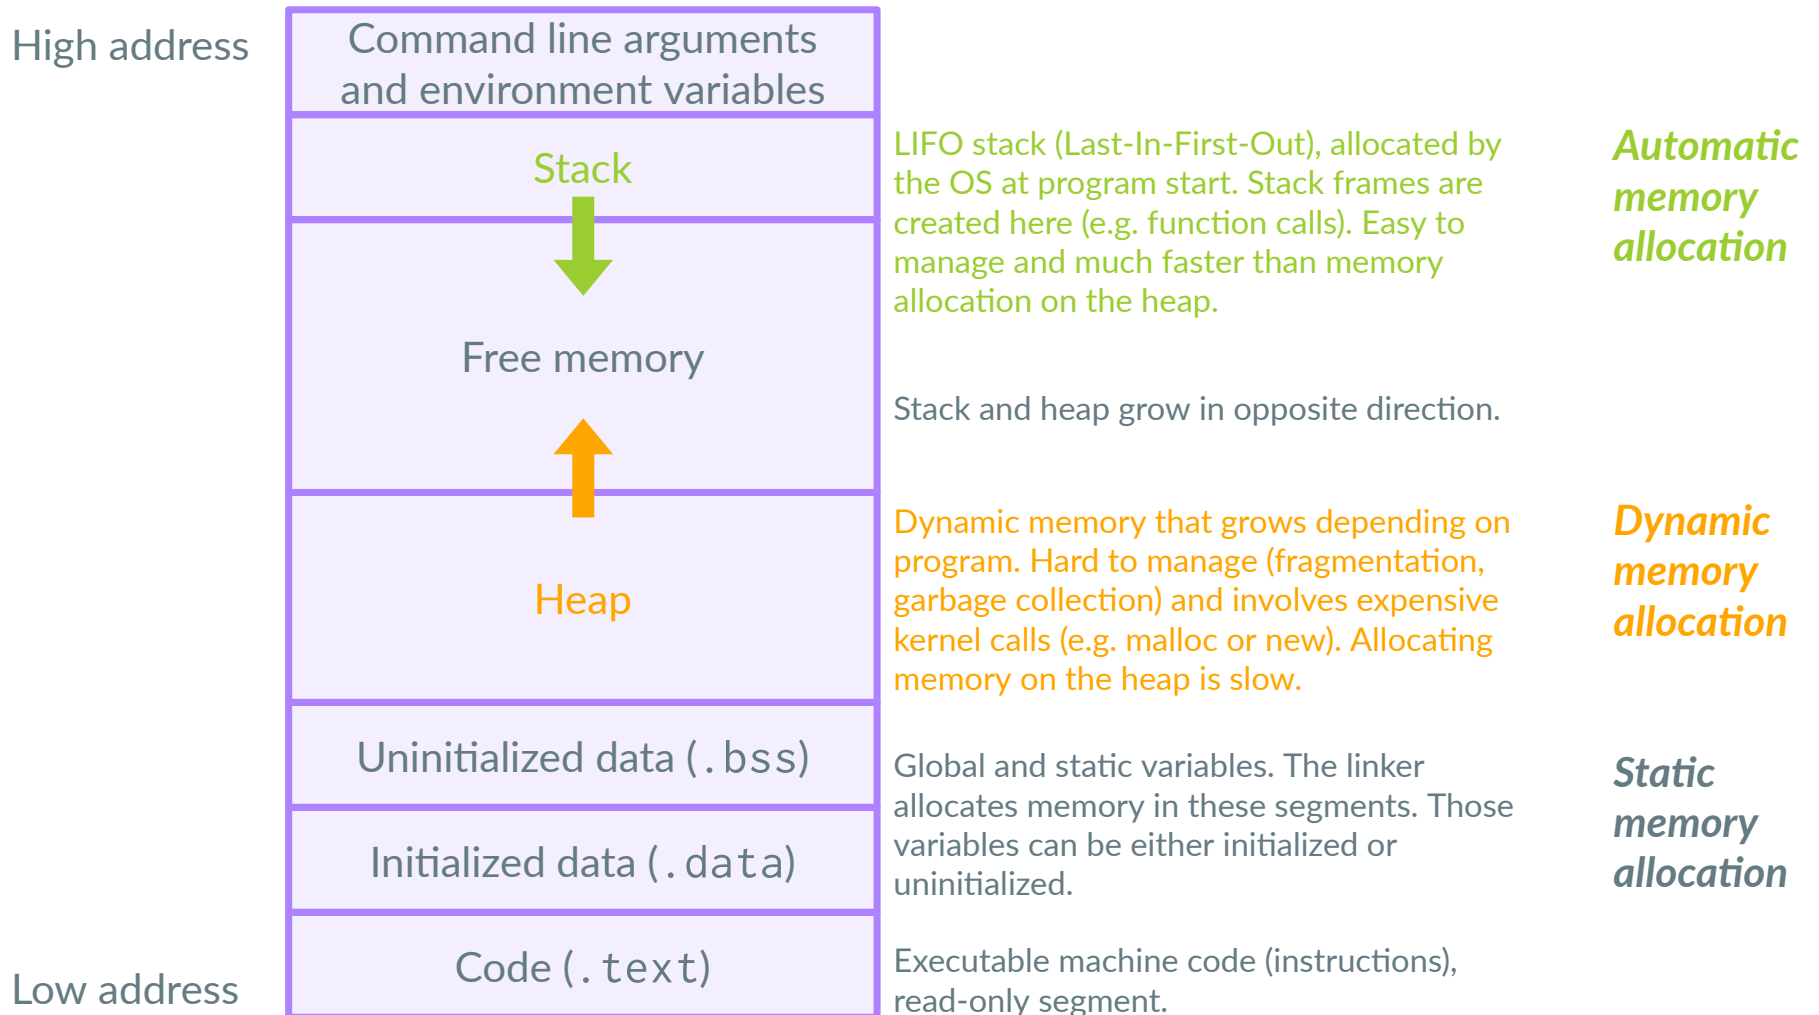
PYTHON INTERNALS: DYNAMIC MEMORY

Different ways (places) to allocate memory:

- **Static memory allocation:** variables with known size at compile time. The compiler allocates this memory inside the *executable*.
- **Automatic memory allocation:** similar to static memory allocation (the allocation requirements are known at compile time). Allocation is carried out on the stack when the code executes (for example a function body).
- **Dynamic memory allocation:** when it is not possible for the compiler to determine a specific memory request the memory must be allocated dynamically. A dynamic memory segment is allocated on the heap.
Example: the allocation size depends on user input.
- All objects in Python (including frame objects and code objects) are **allocated dynamically on the heap** → a reason why Python is slow!

PYTHON INTERNALS: DYNAMIC MEMORY

Virtual memory of a Linux process:



PYTHON INTERNALS: DYNAMIC MEMORY

- The Python interpreter *emulates* a frame stack using *dynamic memory*. Frame objects are *pushed* and *popped* to and from the frame stack on the *heap* (dynamic memory pool).
- These "stack" operations are more expensive than true operations on a stack with *automatic memory allocation* as it is the case for a x86_64 executable for example.
- Since all Python objects are allocated on the heap, *generator objects persist* until the interpreter explicitly removes them from the heap. This allows to easily resume a suspended generator including its state. Because a PyFrameObject stores the last instruction in *f_lasti*, it will be used to index into the bytecode of a generator object to resume execution with instruction $f_lasti + 1 \rightarrow$ that is, after the last active yield statement.

PYTHON INTERNALS: SUMMARY

- The Python interpreter exposes a number of internal objects to the user of which we have discussed three:
 - Code objects for byte compiled code
 - Frame objects to execute code.
 - Generator objects for suspension and resumption of code execution.
 - Traceback objects for debugging (not discussed).
- It is rare that you will need to manipulate these objects directly in your Python code.
- We have discussed them here to understand the low-level Python internals without going too deep into the interpreter source code.
- The most important source code of the CPython interpreter is `ceval.c`.

PYTHON LIST OBJECTS AND NUMPY ARRAYS

- We have seen that excessive use of dynamic memory allocation in Python is a cause for its reputation of being slow.
- It is not due to bad design (the GIL is debatable), but rather the cost of prioritizing *flexibility* and the possibility for *fast prototyping* which has value in its own right.
- One of the reasons for this performance penalty is that Python objects are not necessarily near by in memory due to dynamic memory allocation and object oriented design.
- *Why does this matter since memory cells in random access memory (RAM) can be accessed in constant time you may ask?*
 - Additional pointer dereferences until you get to the data → *in Python everything is an object and must be referenced by pointers.*
 - *Spatial and temporal locality* of the data is not optimal. Results in many cache misses when reading or writing data.

PYTHON LIST OBJECTS AND NUMPY ARRAYS

Python built-in lists:

- Let us see how list objects are implemented in Python:

```
1 typedef struct {
2     PyVarObject ob_base;
3     // Vector of pointers to
4     // list elements. list[0]
5     // is ob_item[0], etc.
6     PyObject **ob_item;
7     Py_ssize_t allocated;
8 } PyListObject;
```

- `ob_item` is a pointer to pointer(s) to `PyObject`'s. **Example:** `ob_item[0]` returns a pointer to a `PyObject`, `ob_item[1]` returns the next pointer to the second `PyObject` and so on.

- In the following we assume `PyObject` represents a Python integer:

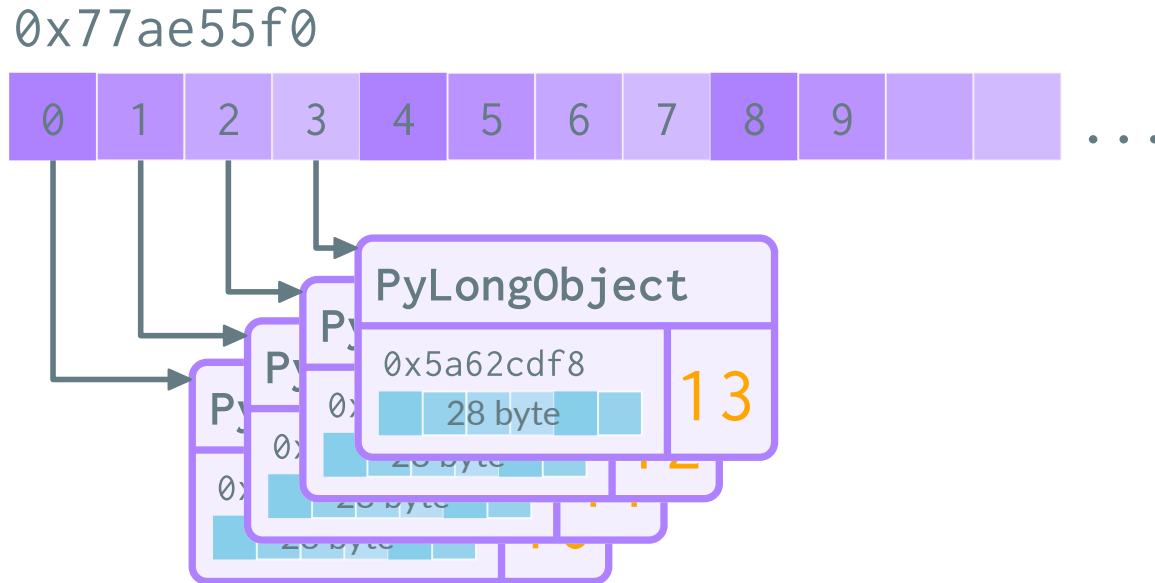
```
1 struct _longobject {
2     PyObject ob_base;
3     Py_ssize_t ob_size;
4     // the actual integer:
5     digit ob_digit[1];
6 } PyLongObject;
```

We assume that the `PyObject` takes **16 byte**, `ob_size` is **8 byte** and `ob_digit` is **4 byte**. A `PyLongObject` then has a size of **28 byte**.

- The actual integer value is stored at `ob_digit[1]`.

PYTHON LIST OBJECTS AND NUMPY ARRAYS

Python built-in lists:



```
1 li = [10, 11, 12, 13, 14, 15, 16, 17, 18, 19] # assume this list of integers
```

- The elements of `ob_item` are *coalesced* in memory. We can access the `Py0bject` references in the list with $\mathcal{O}(1)$ complexity.
- To obtain the actual value we must *dereference* the pointer and read `ob_digit[0]` for every item in the list!

PYTHON LIST OBJECTS AND NUMPY ARRAYS

We can visualize this list on pythontutor.com showing all heap allocations:

The screenshot displays the Python Tutor interface for Python 3.6. The code editor shows a single line: `1 li = [10, 11, 12, 13, 14, 15, 16, 17, 18,]`. A red arrow points to the line number '1', and a green arrow points to the code. Below the code, a legend indicates: a green arrow for 'line that just executed' and a red arrow for 'next line to execute'. A horizontal scrollbar is positioned below the legend. At the bottom of the code editor, there are two buttons: '< Prev' and 'Next >'. Below these buttons, it says 'Step 1 of 1'. At the very bottom, it says 'Visualized using [Python Tutor](#)' and '[Customize visualization](#)'. To the right of the code editor, there are two empty panels labeled 'Frames' and 'Objects'.

Python 3.6

```
1 li = [10, 11, 12, 13, 14, 15, 16, 17, 18,]
```

[Edit this code](#)

→ line that just executed
→ next line to execute

< Prev Next >

Step 1 of 1

Visualized using [Python Tutor](#)
[Customize visualization](#)

Frames Objects

PYTHON LIST OBJECTS AND NUMPY ARRAYS

Example: for-loop over iterable and sum values

Sum values in iterable x:

```
1 import timeit
2 import numpy as np
3
4 def pysum(x):
5     s = 0
6     for i in x:
7         s += i
8
9 if __name__ == "__main__":
10     x = np.array(list(range(1000000)))
11     t = timeit.timeit('f(x)',
12                       globals={'f': pysum, 'x': x},
13                       number=10)
```

Instructions for pysum code object:

1	5	0	LOAD_CONST	1 (0)
2		2	STORE_FAST	1 (s)
3				
4	6	4	LOAD_FAST	0 (x)
5		6	GET_ITER	
6	>>	8	FOR_ITER	12 (to 22)
7		10	STORE_FAST	2 (i)
8				
9	7	12	LOAD_FAST	1 (s)
10		14	LOAD_FAST	2 (i)
11		16	INPLACE_ADD	
12		18	STORE_FAST	1 (s)
13		20	JUMP_ABSOLUTE	8
14	>>	22	LOAD_CONST	0 (None)
15		24	RETURN_VALUE	

Running this code with 1'000'000 elements takes **0.74 seconds**, averaged over 10 samples.

PYTHON LIST OBJECTS AND NUMPY ARRAYS

- While PyObject's in Python are very flexible (Python is dynamically typed), this flexibility comes at a performance price.
- The Python interpreter is designed to work with PyObject's exclusively → *everything in Python is an object*.
- Performance oriented designs are centered around *data* rather than objects.
- Because the Python interpreter is written in C, extensions can easily be implemented.
- **NumPy** is a Python extension module designed for efficient numerical computation in Python.

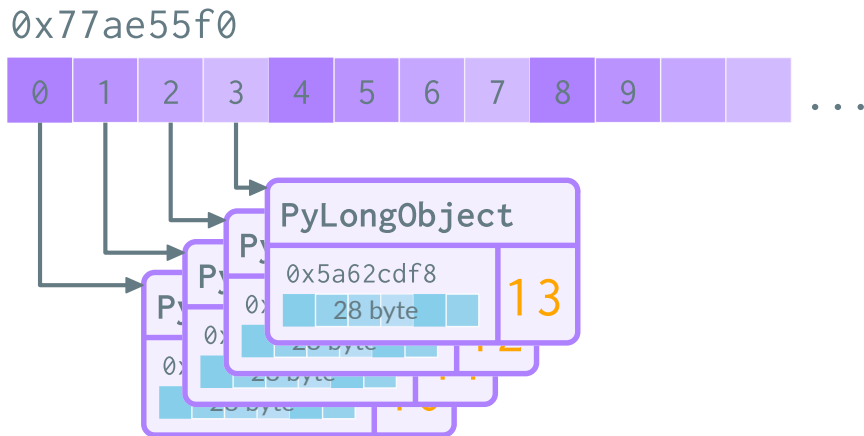
It operates on its own data structures in order not to pay the performance price for flexibility → `np.array`

PYTHON LIST OBJECTS AND NUMPY ARRAYS

Back to our list:

```
1 li = [10, 11, 12, 13, 14, 15, 16, 17, 18, 19] # assume this list of integers
```

Object oriented Python list:




- The list above is the same as

```
1 // 10 contiguous pointer
2 PyObject *ob_item[10];
```

- PyLongObject uses 32-bit integral type for integers, i.e., `int` in C.

Data oriented NumPy array:



- The elements of a NumPy array are **contiguous data items**, not pointers (references) to PyObject's.
- The NumPy array above is similar to

```
1 // 10 contiguous data items
2 int ob_item[10];
```

- Reading the data in this format will saturate the memory bandwidth!

PYTHON LIST OBJECTS AND NUMPY ARRAYS

Example: for-loop over iterable and sum values (as before)

Sum values in iterable x:

```
1 import timeit
2 import numpy as np
3
4 def pysum(x):
5     s = 0
6     for i in x:
7         s += i
8
9 def npsum(x):
10     s = x.sum()
11
12 if __name__ == "__main__":
13     x = np.array(list(range(1000000)))
14     t = timeit.timeit('f(x)',
15                       globals={'f': npsum, 'x': x},
16                       number=10)
```

Instructions for npsum code object:

1	10	0	LOAD_FAST	0 (x)
2		2	LOAD_METHOD	0 (sum)
3		4	CALL_METHOD	0
4		6	STORE_FAST	1 (s)
5		8	LOAD_CONST	0 (None)
6		10	RETURN_VALUE	

Running this code with 1'000'000 elements averaged over 10 samples:

- Pure Python: **0.74 seconds**
- NumPy array: **0.0046 seconds**

→ *Two orders of magnitude faster!*

Note: the `sum()` built-in function is only slightly faster (**0.60 seconds**) than the naive for-loop implementation!

RECAP

- Python internals:
 - Code objects and bytecode
 - The interpreter and the evaluation loop
 - Frame objects
 - Generator objects
- Why are Python built-in lists slow and NumPy arrays fast?

Further reading:

- Python bytecode disassembler: <https://docs.python.org/3/library/dis.html>
- CPython interpreter source code: <https://github.com/python/cpython>
- Python/C API Reference Manual: <https://docs.python.org/3/c-api/index.html>
- Chapters "The Evaluation Loop", "Memory Management" and "Objects and Types" in Anthony Shaw, *CPython Internals: Your Guide to the Python 3 Interpreter*, Real Python, 2020